

Assignment No. 3 (Due: April 18, 2017)

Score: ____/100

Student: Hualiang Li hualiang@hawaii.edu

Date: April 10, 2017

31 (20/100). Assume a computer architecture with a 1024-entry Branch History Table (BHT). Each BHT entry is a two bit predictor that implements a two-bit saturating counter that starts in the strong not-taken state. After the following code execution, what is the state of the BHT? What is the prediction accuracy of each branch?

```

p_4:
// Assume that at address 0x1000, there is an array of word sized integers
// with the following data [ 0x0, 0x0, 0x0, 0x0, 0x0, 0x1, 0x0, 0x1 ]
ADDI R7, R0, 7
ADDI R11, 0x1000
loop:
SUBI R7, R7, 1
ANDI R8, R7, 1
b_0:
BEQZ R8, l_1
LW R12, 0(R11)
b_1:
BEQZ R12, l_2
l_1:
ADD R9, R9, R16
l_2:
ADDI R11, R11, 4
b_3:
BNEZ R7, loop

```

Iteration	b_0 state	b_1 state	b_3 state
1	Taken NT	SNT	Taken NT
2	Ntaken SNT	Taken NT	Taken T
3	Taken NT	NT	Taken ST
4	Ntaken SNT	Taken T	Taken ST
5	Taken NT	T	Taken ST
6	Ntaken SNT	Ntaken NT	Taken ST
7	Taken NT	NT	Ntaken T

For b_0, the accuracy is 42.9%. For b_1, the accuracy is 0%. For b_3, the accuracy is 57.1%

32 (20/100). [PH11], page 137, problem 2.9-a,b,c,d.

- From 2.8 c, the memory access time is 1.69ns for 4-way 64KB. The direct mapped cache access time is 0.86ns. $0.0033*20+(0.8*2+(1-0.8)*3)*(1-0.0033) = 2.3$ cycles. So, $2.3*0.5 = 1.15$ ns
- $0.83/0.5 = 1.66$. So 66% faster.
- $0.0033*20+(0.8*2+(1-0.8)*15)*(1-0.0033) = 4.65$ cycles. So $4.65*0.5 = 2.33$ ns
- $2.4/1.59 = 1.509$. So 50.9% faster.

33 (20/100). You are executing on a VMIPS (as described on page 266 of [PH11]) architecture the code below. Assume that the maximum vector length of the architecture is 128. Draw the pipeline diagram of this code executing on a single-lane architecture which has an independent load unit, store unit, multiply unit, and ALU unit. Loads have a latency of three cycles (L0, L1, L2), stores take two cycles to occur (S0, S1), multiplies take 5 cycles (Y0, Y1, Y2, Y3, Y4), and ALU operations take two cycles (X0, X1). Assume full pipelining of the functional units. Assume that the pipeline has a dedicated register read stage and a single write-stage. For the first part of this problem, assume that the architecture support chaining through the register file, but only has two read ports and one write port on the register file.

```
C Code :
for (i = 0; i < 6; i++)
{
a[i] = (b[i] * 7) + c[i];
}
```

VMIPS:

ADDI R7, R0, 7

CVT.W.D F2, R7

ADDI R1, R0, 6

MTC1 VLR, R1

LV V1, R4

MULVS.D V3, V1,

F2 LV V2, R5

ADD V4, V3, V2

SV R6, V4

LV F D R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

MULVS.D F D D D D D R Y0 Y1 Y2 Y3 Y4 W

R Y0 Y1 Y2 Y3 Y4 W

R Y0 Y1 Y2 Y3 Y4 W

R Y0 Y1 Y2 Y3 Y4 W

R Y0 Y1 Y2 Y3 Y4 W

R Y0 Y1 Y2 Y3 Y4 W

LV F F F F F F D D D D D D D D R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

R L0 L1 L2 W

ADD F F F F F F F F D D D D D D R X0 X1 W

R X0 X1 W

R X0 X1 W

R X0 X1 W

R X0 X1 W

R X0 X1 W

SV F F F F F F F D D D D D D R S0 S1 W

R S0 S1 W

R S0 S1 W

R S0 S1 W

R S0 S1 W

R S0 S1 W

34 (20/100). [PH11], page 254, 3.13.

Processor A:

Time

1 T1 LD R1 T2 LD R1

2 T3 LD R1 T4 LD R1

3 T1 LD R2 T2 LD R2

4 T3 LD R2 T4 LD R2

.

.

.

16 T3 LD R8 T4 LD R8

17 T1 BEQ T2 BEQ

.

.

.

48 T3 BEQ T3 BEQ

51	T1	BCT	T2	BLT
----	----	-----	----	-----

.

.

.

54

So, for processor A, $54 * 2 = 108$ cycles

Processor B:

Time

1	T1	LDR1	LDR2	LDR3	LDR4
2	T1	LDR5	LDR6	LDR7	LDR8
3	T1	BEQ			
4	T2	LDR1	LDR2	LDR3	LDR4
5	T2	LDR5	LDR6	LDR7	LDR8
6	T2	BEQ			

.

.

.

10	T3	LDR1	LDR2	LDR3	LDR4
----	----	------	------	------	------

11	T3	LDR5	LDR6	LDR7	LDR8
----	----	------	------	------	------

12	T3	BEQ			
----	----	-----	--	--	--

13	T1	BEQ			
----	----	-----	--	--	--

14	T2	BEQ			
----	----	-----	--	--	--

15	T3	BEQ			
----	----	-----	--	--	--

16	T4	BEQ			
----	----	-----	--	--	--

.

.

.

37	T1	BEQ			
----	----	-----	--	--	--

38	T2	BEQ			
----	----	-----	--	--	--

39	T3	BEQ			
----	----	-----	--	--	--

40	T4	BEQ			
----	----	-----	--	--	--

41	T1	DADDIU			
----	----	--------	--	--	--

42	T2	DADDIU			
----	----	--------	--	--	--

43	T3	DADDIU			
----	----	--------	--	--	--

44	T4	DADDIU			
----	----	--------	--	--	--

45	T1	BLT			
----	----	-----	--	--	--

46	T2	BLT			
----	----	-----	--	--	--

47	T3	BLT			
----	----	-----	--	--	--

48	T4	BLT			
----	----	-----	--	--	--

So, for processor B, $48 * 2 = 96$ cycles

Processor C

Time

1	T1	LDR1	LDR2	LDR3	LDR4	LDR5	LDR6	LDR7	LDR8
---	----	------	------	------	------	------	------	------	------

2	T2	LDR1	LDR2	LDR3	LDR4	LDR5	LDR6	LDR7	LDR8
---	----	------	------	------	------	------	------	------	------

3	T3	LDR1	LDR2	LDR3	LDR4	LDR5	LDR6	LDR7	LDR8
---	----	------	------	------	------	------	------	------	------

4	T4	LDR1	LDR2	LDR3	LDR4	LDR5	LDR6	LDR7	LDR8
---	----	------	------	------	------	------	------	------	------

5	T1	BEQ							
---	----	-----	--	--	--	--	--	--	--

6	T2	BEQ							
---	----	-----	--	--	--	--	--	--	--

7	T3	BEQ							
---	----	-----	--	--	--	--	--	--	--

8	T4	BEQ							
---	----	-----	--	--	--	--	--	--	--

.

.

.

36	T4	BEQ							
----	----	-----	--	--	--	--	--	--	--

37	T1	DADDIU							
----	----	--------	--	--	--	--	--	--	--

38	T1	BLT							
----	----	-----	--	--	--	--	--	--	--

39	T2	BLT							
----	----	-----	--	--	--	--	--	--	--

40	T3	BLT							
----	----	-----	--	--	--	--	--	--	--

41	T4	BLT							
----	----	-----	--	--	--	--	--	--	--

So, for processor c, $41 * 2 + 1 = 83$ cycles, $83 * 4 = 332$ cycles.

35 (20/100). You are writing a multi-threaded program that will count the number of occurrences of a value in an array. The values in the array are between 0 and 1023. In effect, you will be building a histogram. Assume that the list of numbers is very large, on the order of gigabytes large. Extend the following program such that 100 threads (processors) can execute on the program concurrently. Assume a sequentially consistent memory model. Add P() and V() semaphores where appropriate and add any storage needed for the semaphores. Explain why the speedup of such a solution may not be 100x. Note that the output lock array is assumed to be initialized to 1 (this allows for a mutex).

```
// Sequential code, assume that the input and output arrays are created
// outside of the function
#define MAX_VALUE 1023
function(int input_array_size, int * input_array, int * output_array)
{
    int counter;
    for(counter = 0; counter < input_array_size; counter++)
    {
        assert(input_array[counter] <= MAX_VALUE);
        assert(input_array[counter] >= 0);
        output_array[input_array[counter]]++;
    }
}
```

```
#include <thread.h>
#include <stdlib.h>
#define MAX 1023
#define MIN 0
#define NUM_THREADS 100

struct data{
    int input_size;
    int* input_array;
    int* output_array;
    int* output_lock_array;
}

struct thread_data thread_data_array[NUM_THREADS];

void function(void * thread_args)
{
    struct thread_data * temp_data;
    temp_data = (struct thread_data *) thread_args;
    int input_array_size = temp_data->input_size;
    int * input_array = temp_data->input_array;
    int * output_array = temp_data->output_array;
    int * output_lock_array = temp_data->output_lock_array;
    int counter;
    for(counter = 0; counter < input_array_size; counter++)
    {
        assert(input_array[counter] <= MAX);
        assert(input_array[counter] >= MIN);
        P(&output_lock_array[counter]);
        output_array[input_array[counter]]++;
        V(&output_lock_array[counter]);
    }
}
```

```

}

int main() {

    int counter;
    pthread_t threads[NUM_THREADS];
    int * input_array;
    input_array = malloc(INPUT_ARRAY_ELEMENTS*sizeof(int));
    int * output_array = malloc((MAX + 1) * sizeof(int));
    int * output_lock_array = malloc((MAX + 1) * sizeof(int));
    for(counter = 0; counter <= MAX; counter++);
    {
        output_array = 0;
        output_lock_array = 1;
    }
    for(counter = 0; counter < 100; counter++)
    {
        thread_data_array[counter].input_array_size = INPUT_ARRAY_ELEMENTS /
            NUM_THREADS;
        thread_data_array[counter].input_array =
            &input_array[INPUT_ARRAY_ELEMENTS / NUM_THREADS * counter];
        thread_data_array[counter].output_array = output_array;
        thread_data_array[counter].output_lock_array = output_lock_array;
        pthread_create(&threads[counter], NULL, function, (void
            *)thread_data_array[counter]);
    }
    for(counter = 0; counter < 100; counter++)
    {
        pthread_join(threads[counter], NULL);
    }
}

```

No. It will be less than 100 faster. This is because of 1) the threads overhead, 2) lock overhead 3) cach miss overhead.

References

- [PH11] David A. Patterson and John L. Hennessy. *Computer Architecture: A Quantitative Approach*. 5th ed. Morgan Kaufmann Publishers Inc., 2011. isbn: 978-0-12-383872-8. url: <http://booksite.elsevier.com/9780123838728>.